

Embedded SQL

csc343, Introduction to Databases

Diane Horton

with examples from Ullman and Widom

Winter 2016



UNIVERSITY OF
TORONTO

Problems with using interactive SQL

- Standard SQL is not “Turing-complete”.
 - E.g., Two profs are “colleagues” if they’ve co-taught a course or share a colleague.
 - We can’t write a query to find all colleagues of a given professor because we have no loops or recursion.
- You can’t control the format of its output.
- And most users shouldn’t be writing SQL queries!
 - You want to run queries that are *based on* user input, not have users writing actual queries.

SQL + a conventional language

- If we can combine SQL with code in a conventional language, we can solve these problems.
- But we have another problem:
 - SQL is based on relations, and conventional languages have no such type.
- It is solved by
 - feeding tuples from SQL to the other language one at a time, and
 - feeding each attribute value into a particular variable.

Approaches

- Three approaches for combining SQL and a general-purpose language:
 - Stored Procedures
 - Statement-level Interface
 - Call-level interface

Three Approaches

I. Stored Procedures

- The SQL standard includes a language for defining “stored procedures”, which can
 - have parameters and a return value,
 - use local variables, ifs, loops, etc.,
 - execute SQL queries.
- Stored procedures can be used in these ways:
 - called from the interpreter,
 - called from SQL queries,
 - called from another stored procedure,
 - be the action that a trigger performs.

Example (just to give you an idea)

- A binary function `BandW(y INT, s CHAR(15))` that returns true iff
 - movie studio `s` produced no movies in year `y`, or
 - produced at least one comedy.
- (Yes, that's an odd name for this function.)
- Reference: Ullman and Widom textbook, chapter 9

Reference: textbook figure 9.1.3

```
CREATE FUNCTION BandW(y INT, s CHAR(15)) RETURNS BOOLEAN
IF NOT EXISTS
    (SELECT *
      FROM Movies
      WHERE year = y AND studioName = s)
THEN RETURN TRUE;
ELSIF 1 <=
    (SELECT COUNT(*)
      FROM Movies
      WHERE year = y AND studioName = s AND
            genre = 'comedy')
THEN RETURN TRUE;
ELSE RETURN FALSE;
END IF;
```


Calling it

- Now we can say things like this:

```
SELECT StudioName  
FROM Studios  
WHERE BandW(2010, StudioName);
```

Not very standard

- The language is called **SQL/PSM** (Persistent Stored Modules).
 - It came into the SQL standard in SQL3, 1999.
 - Reference: textbook, section 9.4
- By then, various commercial DBMSs had already defined their own proprietary languages for stored procedures
 - They have generally stuck to them.
- PostgreSQL has defined **PL/pgSQL**.
 - It supports some, but not all, of SQL/PSM.
 - Reference: Chapter 39 of the PostgreSQL documentation.

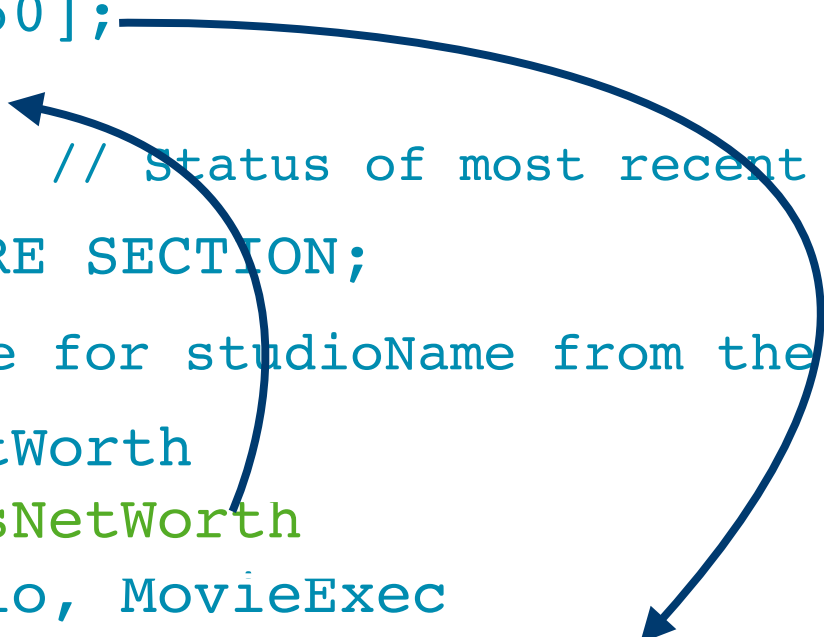
2. Statement-level interface (SLI)

- Embed SQL statements into code in a conventional language like C or Java.
- Use a preprocessor to replace the SQL with calls written in the host language to functions defined in an SQL library.
- Special syntax indicates which bits of code the preprocessor needs to convert.

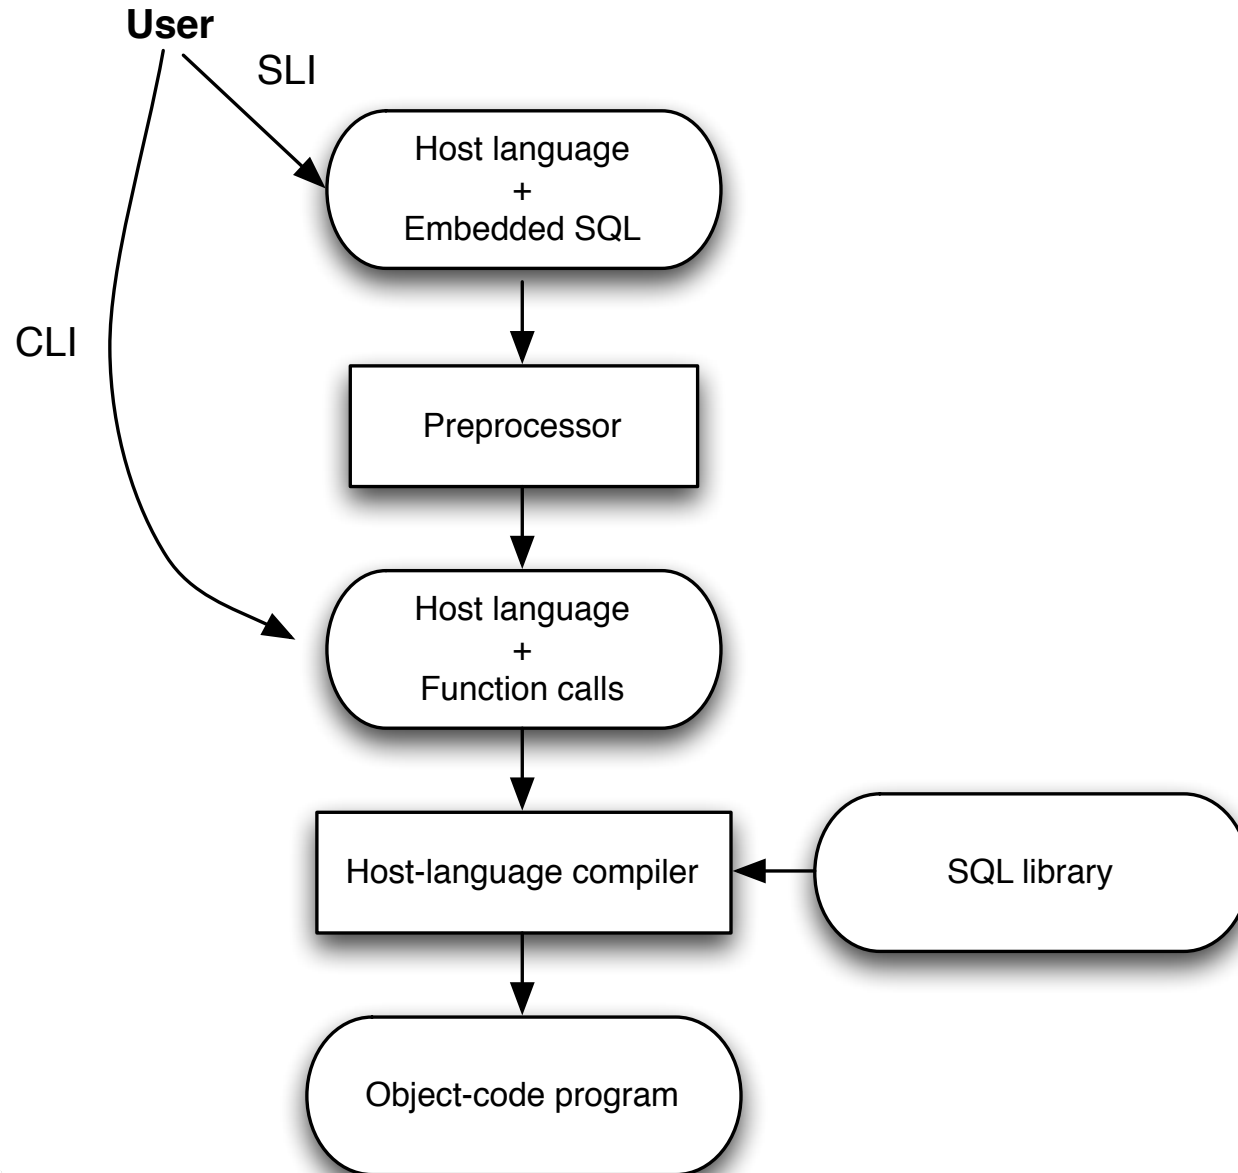
Example (just to give you an idea)

Reference: textbook example 9.7

```
void printNetWorth() {  
    EXEC SQL BEGIN DECLARE SECTION;  
        char studioName[50];  
        int presNetWorth;  
        char SQLSTATE[6]; // Status of most recent SQL stmt  
    EXEC SQL END DECLARE SECTION;  
    /* OMITTED: Get value for studioName from the user. */  
    EXEC SQL SELECT netWorth  
        INTO :presNetWorth  
        FROM Studio, MovieExec  
        WHERE Studio.name = :studioName;  
    /* OMITTED: Report back to the user */  
}
```



Big picture (figure 9.5)



3. Call-level interface (CLI)

- Instead of using a pre-processor to replace embedded SQL with calls to library functions, write those calls yourself.
- Eliminates need to preprocess.
- Each language has its own set of library functions for this.
 - for C, it's called SQL/CLI
 - for Java, it's called JDBC
 - for PHP, it's called PEAR DB
- We'll look at just one: JDBC.

JDBC

Using JDBC on cdf

- You need to run your JDBC code on dbsrv1.
- The PostgreSQL driver for JDBC is on cdf here:
`/local/packages/jdbc-postgresql`
You'll also find an example program and a how-to in that directory.
- To run JDBC code, you need this driver in your classpath.
- Example: Suppose you have a class called Jelly.java.
`javac Jelly.java`
`java -cp ~/bin/postgresql-8.3-607.jdbc4.jar:`
`Jelly`

JDBC Example (see section 9.6)

Do this once in your program:

```
/* Get ready to execute queries. */  
import java.sql.*;  
/* A static method of the Class class. It loads the  
   specified driver */  
Class.forName("org.postgresql.jdbc.Driver");  
Connection conn = DriverManager.getConnection(  
    jdbc:postgresql://localhost:5432/csc343h-dianeh,  
    dianeh,  
    "" );  
/* Continued ... */
```

The arguments to getConnection

- `jdbc:postgresql`

We'll use this, but it could be, e.g., `jdbc:mysql`

- `localhost:5432`

You must use exactly this for cdf.

- `csc343h-dianeh` and `dianeh`

Substitute your cdf userid.

- `""`

Password (unrelated to your cdf password).

Literally use the empty string.

Do this once per query in your program:

```
/* Execute a query and iterate through the resulting
   tuples. */

PreparedStatement execStat = conn.prepareStatement(
    "SELECT netWorth FROM MovieExec");

ResultSet worths = execStat.executeQuery();
while (worths.next()) {
    int worth = worths.getInt(1);
    /* If the tuple also had a float and another int
       attribute, you'd get them by calling
       worths.getFloat(2) and worths.getInt(3).
       Or you can look up values by attribute name.
       Example: worths.getInt(netWorth)
    */

    /* OMITTED: Process this net worth */
}
```

Exceptions can occur

- Any of these calls can generate an exception.
- Therefore, they should be inside try/catch blocks.

```
try {  
    /* OMITTED: JDBC code */  
} catch (SQLException ex) {  
    /* OMITTED: Handle the exception */  
}
```

- The class `SQLException` has methods to return the `SQLSTATE`, etc.

What is “preparation”?

- Preparing a statement includes parsing the SQL, compiling and optimizing it.
- The resulting `PreparedStatement` can be executed any number of times without having to repeat these steps.

If the query isn't known until run time

- You may need input and computation to determine the query.
- You can hard-code in the parts you know, and use “?” as a placeholder for the values you don't know.
- This is enough to allow a `PreparedStatement` to be constructed.
- Once you know values for the placeholders, methods `setString`, `setInt`, etc. let you fill in those values.

Example (figure 9.22)

```
PreparedStatement studioStat =  
    conn.prepareStatement(  
        "INSERT INTO Studio(name, address)  
        VALUES(?, ?)"  
    );
```

```
/* OMITTED: Get values for studioName and studioAddr */  
studioStat.setString(1, studioName);  
studioStat.setString(2, studioAddr);  
studioStat.executeUpdate();
```

Why not just build the query in a string?

- We constructed an incomplete `PreparedStatement` and filled in the missing values using method calls.
- Instead, we could just build up the query in an ordinary string at run time, and ask to execute that.
- There are classes and methods that will do this in JDBC.

Example that builds the query in a string

- We can just use a `Statement`, and give it a `String` to execute.

```
Statement stat = conn.createStatement();
String query =
    "SELECT networth
    FROM MovieExec
    WHERE execName like '%Spielberg%';
"
ResultSet worths = stat.executeQuery(query);
```

What could possibly go wrong?

Queries vs updates in JDBC

- The previous examples used `executeQuery`.
- This method is only for pure queries.
- For SQL statements that change the database (insert, delete or modify tuples, or change the schema), use the analogous method `executeUpdate`.